

Evidence Helps on Average, but Sometimes Hurts: Characterizing Evidence-Induced Regressions in LLM Fact Verification

Chimdumebi Mitchell Nebolisa
Affiliation TODO

Jessica Udry
Affiliation TODO

Corrected submission draft for coauthor review (2026-09-28); venue formatting TODO

Abstract

Supplying designated gold evidence is widely assumed to make language-model fact verification more reliable. We test that assumption in a paired experiment on 10,000 balanced FEVER Supported/Refuted claims, evaluating GPT-5.4 and GPT-5.4-mini without evidence and with FEVER-designated gold evidence (40,000 predictions). Evidence raises aggregate accuracy from 88.69% to 96.04% (+7.35 pp) for GPT-5.4 and from 84.67% to 95.54% (+10.87 pp) for GPT-5.4-mini. Those gains conceal 226 unique claims that were correct without evidence and become wrong after the same gold evidence is supplied (40 shared by both models, 186 specific to one). We describe these regressions with blinded five-judge progressive-disclosure adjudication by two model families, Grok and Claude, after repairing a structured-evidence stage whose historical inputs displayed corrupted sentence text. A descriptive taxonomy keeps judge decisiveness and agreement with FEVER separate. On the repaired inputs, the panels reach the FEVER label from the evidence sentences alone for 34.1% (Grok) and 46.9% (Claude) of regressions, reach a decisive verdict opposite to FEVER for 19.5% and 23.9%, and remain nondecisive for 31.4% and 18.1%. Claude is more decisive than Grok at every disclosure stage; titles and structured evidence resolve a minority of nondecisive cases and never switch a decisive verdict to the opposite label; and exact per-claim category agreement between families is 76.1% ($\kappa = 0.67$). The structured-stage analysis follows a dated amendment, written after judging, that analyses each job’s accepted output; the results are stable when earlier valid outputs are substituted. Evidence can raise overall accuracy while creating localized failures, and how those failures distribute across adjudication outcomes depends on the adjudicating model family; the outcomes do not identify why GPT changed its answers.

Keywords: fact verification, FEVER, large language models, silver adjudication, LLM-as-judge, evidence-induced regressions

1 Introduction

Evidence grounding is a standard method for making language-model fact verification more trustworthy. If a model is shown the evidence that a benchmark designates as sufficient, aggregate accuracy should rise, and the errors that remain should reflect problems of retrieval or coverage rather than a failure to use that evidence. Aggregate accuracy can, however, hide a second

phenomenon: claims that a model classifies correctly without evidence but incorrectly after the designated evidence is added. We call such a correct-to-wrong change a *regression*.

We study regressions on FEVER [Thorne et al., 2018a], a claim-verification dataset built from Wikipedia. Annotators labeled each FEVER claim Supported, Refuted or Not Enough Info and, for Supported and Refuted claims, recorded the Wikipedia sentences that serve as evidence. FEVER suits the question because this annotated evidence can be given to a model directly. From its development set we draw a balanced sample of 10,000 Supported and Refuted claims and evaluate GPT-5.4 and GPT-5.4-mini on each claim twice: once with the claim alone, and once with the claim plus its FEVER-designated evidence, which we call gold evidence. Because each claim is compared with itself and no retrieval step is involved, a change between the two conditions can be attributed to the supplied evidence rather than to retrieval. The sample is restricted to the two labels for which FEVER records evidence, so it does not represent FEVER as a whole.

Evidence helps in aggregate (Table 1), but it also produces regressions. For GPT-5.4, 115 of the 10,000 claims (1.15%) changed from correct without evidence to wrong with it; for GPT-5.4-mini, 151 (1.51%) did; and 226 distinct claims regressed for at least one model. These regressions are rare, but they are not unstructured noise. They concentrate on claim–evidence pairs that local natural language inference (NLI) models judge weakly supported by the evidence, and only 40 of the 226 claims regress for both models.

A regression is an observed change in a prediction; it does not reveal its cause. The same change could arise because the model did not use evidence that is in fact decisive, because the evidence can reasonably be read in more than one way, or because the evidence does not settle the claim even for a careful reader. The original runs recorded only predicted labels, so they cannot distinguish these possibilities. We therefore ask an observable question instead: how do separate, blinded readers judge the same evidence? To answer it we use progressive-disclosure silver adjudication. “Silver” means that the verdicts come from automated model judges rather than expert human annotators, so they are not ground truth. “Progressive disclosure” means that each judge sees the evidence in three cumulative stages: the selected evidence sentences, exactly as GPT received them (Stage A); the same sentences with their Wikipedia page titles (Stage B); and structured FEVER evidence with annotation boundaries and any alternative evidence sets (Stage C). Recording verdicts at every stage shows whether a claim is decided from the sentences GPT saw, only after more context, or not at all. At each stage, five isolated runs of one locked model judge every claim, and a descriptive taxonomy records, separately, whether their consensus is decisive (Supported or Refuted), whether a decisive verdict agrees with FEVER, and the stage at which agreement first appears. Two model families serve as judges: a locked Grok model judges a 1,060-claim diagnostic cohort that includes the 226 regressions, and a locked Claude model judges the 226 regressions.

An audit of the historical adjudication found that its Stage C packets displayed corrupted sentence text for most claims and that its three-way failure-mode taxonomy conflated disagreement with FEVER with evidence utilization. Under a post hoc correction protocol frozen before any new judgment, Stage C was rebuilt from archival FEVER data and re-judged by both families, while the historical Stage A and B judgments, whose inputs were verified unchanged, are retained. Because the protocol’s execution rules could not be followed as written, the analysis follows a dated post-judgment amendment that fixes which output of each judging job is analysed, and we report how sensitive the results are to that choice (Section 3.9).

The study yields two findings. First, designated gold evidence improves overall accuracy while

producing a measurable set of localized regressions. Second, both judge panels describe these regressions as a mixture of FEVER-agreeing, FEVER-disagreeing and nondecisive outcomes, but they assign individual claims to different categories often enough that the proportions depend on which model family judges. We treat that dependence as a property of the measurement instrument: an automated description of the regressions can be relied on only as far as it holds across judge families. Neither finding reveals why GPT changed its answers.

Contributions. We (i) quantify paired evidence gains and evidence-induced regressions on a verified 10,000-claim FEVER experiment; (ii) characterize regressions with confirmatory statistics, two local NLI diagnostics, and a 1,061-claim diagnostic cohort; (iii) describe the 226 unique regressions with blinded progressive-disclosure adjudication on repaired evidence, keeping decisiveness and FEVER agreement separate; (iv) compare two adjudicating families on the same claims and identify which descriptive findings depend on the family; and (v) document an evidence-reconstruction defect in the historical adjudication and its correction. These aims correspond to the four research questions of Section 3.1: whether gold evidence improves aggregate accuracy (RQ1), how frequent and how consistent across models the regressions are (RQ2), which observable adjudication outcomes characterize them (RQ3), and which of those outcomes depend on the judge family (RQ4).

2 Related Work

2.1 Evidence-based fact verification

Automated fact-checking is commonly organized into claim detection, evidence retrieval, and claim verification; the last stage includes verdict prediction and, where applicable, justification production [Guo et al., 2022]. FEVER [Thorne et al., 2018a] makes evidence retrieval and verdict prediction measurable at scale. Its 185,445 claims were generated by altering sentences extracted from Wikipedia, each claim is labeled Supported, Refuted or NotEnoughInfo, and for Supported and Refuted claims annotators recorded the Wikipedia sentences that serve as evidence. The 2018 FEVER shared task, the first held on the dataset, asked systems to retrieve such evidence from Wikipedia and classify each claim [Thorne et al., 2018b]. Its best system reached a FEVER score of 64.21%, a score that credits a label only when the retrieved evidence is also correct, which shows that even with a Wikipedia corpus, finding the right evidence and reaching the right verdict remained difficult. Later datasets extended claim verification beyond FEVER’s setting: SciFact pairs expert-written scientific claims with research abstracts that support or refute them [Wadden et al., 2020]; VitaminC builds contrastive claim–evidence pairs from Wikipedia revisions, in which nearly identical evidence supports a claim in one version and not in the other [Schuster et al., 2021]; HoVer requires evidence drawn from up to four Wikipedia articles [Jiang et al., 2020]; and FEVEROUS annotates evidence as Wikipedia sentences, table cells, or both [Aly et al., 2021].

Benchmark evidence does not guarantee that a model’s verdict depends on it. On FEVER, classifiers that see only the claim perform competitively with evidence-aware models by exploiting cues in the claim’s wording [Schuster et al., 2019]. A model’s accuracy when evidence is supplied therefore does not, by itself, show what the evidence contributed; that requires comparing the same claims with and without evidence.

2.2 LLM reasoning with external evidence

Supplying evidence to a large language model (LLM) does not ensure that the model uses it as intended. In question answering, models can fall back on memorized knowledge when a supplied passage contradicts it [Longpre et al., 2021]. How well they use relevant information in a long input depends on where that information appears [Liu et al., 2024], and irrelevant material in the prompt can divert their reasoning [Shi et al., 2023]. When sources conflict, models weigh a document’s relevance to the query heavily while largely ignoring stylistic features, such as scientific references or a neutral tone, that humans consider important [Wan et al., 2024]. For fact verification on FEVER specifically, Mamta and Cocarascu [2025] find that LLM verdicts are brittle under small perturbations of the input. Together, these studies show that evidence in the input can be ignored, outweighed or misread. They do not measure how often adding correct, benchmark-designated evidence turns a correct answer into a wrong one.

2.3 Evidence sufficiency and ambiguity

Evidence that appears in a prompt is not necessarily evidence that a reader finds sufficient. Atanasova et al. [2022] remove parts of the evidence supplied to fact-checking models and find that the models often fail to recognize when what remains is no longer enough to reach a verdict. Glockner et al. [2024] document claim–evidence pairs that admit conflicting but valid interpretations and argue against forcing each claim into a single label. For a study that asks readers to judge designated evidence, this means that each verdict must be interpreted within its limits. A reader who reaches no decisive verdict has not shown that the evidence is insufficient, and a reader whose decisive verdict opposes the benchmark label has disagreed with that label, not shown it to be wrong.

2.4 LLMs as judges

An LLM judge is a language model prompted to evaluate an input against a rubric in place of a human rater. LLM judges make large-scale evaluation feasible but bring their own biases. Zheng et al. [2023] document position, verbosity and self-enhancement biases; Liu et al. [2023] raise the concern that LLM evaluators favor LLM-generated text; and Chen et al. [2024] find both human and LLM judges vulnerable to misinformation-oversight, gender, authority and beauty biases. Position bias varies across judges and tasks rather than acting as a uniform artifact [Shi et al., 2025], and fine-tuned judge models are not a general substitute for GPT-4 as a judge [Huang et al., 2025]. A description produced by one judge model therefore cannot be assumed to transfer to a judge from another model family.

2.5 Gap addressed by this study

Aggregate accuracy gains and existing research on evidence use leave two questions open. The first is how often supplying correct, benchmark-designated evidence changes a correct answer into a wrong one. Condition-level accuracy cannot show this, because gains on some claims offset losses on others. The second is whether an automated description of those changes is stable when a judge from a different model family applies the same rubric.

Our experiment addresses the first question by supplying FEVER-designated gold evidence, the annotated evidence placed directly in the model input, and comparing each claim with and without it. Because no retrieval step is involved, retrieval failure cannot explain a regression within this paired experiment. Gold evidence does not, however, guarantee that the model attends to it, that every reader finds it sufficient, or that the internal cause of a regression is known. The cost of the design is scope: it studies 2017-era Wikipedia claims restricted to FEVER’s Supported and Refuted labels, with evidence that is designated rather than retrieved.

We address the second question with two judge panels, one from the Grok family and one from the Claude family. Each panel consists of five isolated runs of a locked model that apply the same rubric to the same evidence; they are model-based judges, not independent human experts. Where the two panels agree, a description of the regressions does not depend on one model family. Where they disagree, the disagreement measures sensitivity to the judge family, and we report both panels rather than treating their two values as bounds on a true prevalence. A panel’s failure to reach a decisive verdict after all disclosure stages is recorded as an outcome in its own right, a model-judge analogue of the interpretive ambiguity documented by Glockner et al. [2024].

3 Study Design

3.1 Research questions

The study has two parts. The first measures what supplying designated evidence does to GPT predictions, both in aggregate and claim by claim, which exposes the regressions that aggregate accuracy hides. The second describes those regressions through blinded judge panels whose readings of the same evidence can be compared across model families. This study seeks to answer four questions, two for each part:

- RQ1** How does providing FEVER-designated gold evidence affect aggregate LLM fact-verification accuracy?
- RQ2** When evidence changes a previously correct prediction into an incorrect one, how frequent and consistent are those regressions across models?
- RQ3** Which observable adjudication outcomes—decisiveness, agreement with FEVER, and the disclosure stage at which they change—characterize evidence-induced regressions under progressive disclosure?
- RQ4** Which of those outcomes depend on the adjudicating model family?

3.2 FEVER dataset and sampling

We sample from the FEVER development set [Thorne et al., 2018a]. The manuscript experiment is a balanced 10,000-claim extract: 5,000 Supported and 5,000 Refuted, seed 42, `NOT ENOUGH INFO` excluded. Gold evidence is resolved from official FEVER Wikipedia shards into the text shown to the models. Sampling and validation artifacts are committed with the frozen repository; third-party FEVER and Wikipedia text are not relicensed by the project MIT license.

3.3 Models and paired evidence conditions

Each claim is evaluated four times: GPT-5.4 and GPT-5.4-mini, each under `claim_only` and `claim_plus_evidence`. That yields 40,000 predictions. The evidence condition places the designated FEVER evidence (gold evidence) directly in the prompt; no retrieval step is involved. Because the two conditions differ only in whether that evidence is present, a change in a claim’s prediction between them reflects the supplied evidence rather than retrieval. Original GPT runs stored predicted labels only, so later analysis cannot recover confidence, rationales, or token-level attributions.

3.4 Transition definitions

For each model, a claim has one of four paired transitions: *robust* (correct without and with evidence), *rescue* (wrong without, correct with), *resistant* (wrong in both), or *regression* (correct without, wrong with). A unique-claim regression is a claim that regresses for at least one of the two GPT models. Unique-claim counts are never treated as model-observations: 115 GPT-5.4 regressions and 151 mini regressions are 226 unique claims, not 266.

3.5 Analysis v2

Analysis v2 treats the 40,000-row experiment file as immutable source data. It rebuilds the paired dataset, runs McNemar tests and bootstrap confidence intervals, extracts linguistic and evidence-structure features, embeds claim–evidence pairs, and applies two local NLI models (`nli-deberta-v3-small` and `distilbert-base-uncased-mnli`) as diagnostics, not as ground truth. Leakage audits and an independent 69-check headline verifier are part of the frozen record. Confirmatory tests were pre-specified; NLI and feature analyses are exploratory.

3.6 Diagnostic cohort

The adjudication cohort contains 1,061 unique claims: all 226 unique regressions, 335 resistant, 250 rescue, and 250 robust. One item (SA-000352) is provider-filtered, so judged $N = 1,060$. The Grok panel judges the 1,060-claim cohort; the Claude panel judges only the 226 regressions. This diagnostic sample is not representative of all 10,000 claims, and cross-family comparisons always pair the same 226 claims.

3.7 Progressive-disclosure adjudication

Each judge is a single isolated run of a locked model that applies a fixed rubric to each claim and its evidence and returns a verdict (Supported, Refuted or Ambiguous) with a rating of evidence sufficiency; code combines the five judges’ outputs into a consensus that is Supported, Refuted, Ambiguous or Unresolved. The design discloses the evidence in cumulative stages so that a panel’s verdicts show whether a claim can be decided from the sentences GPT saw, is decided only after more context, or remains undecided. Five isolated judges per stage and batch see: Stage A, the original selected gold-evidence sentence text (identical to the GPT evidence prompt); Stage B, Stage A plus the page titles; Stage C, structured FEVER evidence—ordered sentence pointers, annotation boundaries, and the text of available alternative evidence sets. Stage C can therefore

add information as well as change representation: on the repaired inputs it adds sentence text beyond the selected sentences for 149 cohort claims, including 41 regressions. Judges do not browse, retrieve new pages, or see FEVER labels, GPT predictions, cohort tags, NLI scores, or other judges. Each stage is frozen before the next stage’s packets are built, and packets carry only current-stage evidence. Consensus is computed by code from the five raw judgments; resolver outputs never replace raw consensus. The Grok panel is locked to `cursor-grok-4.6-high-fast` and the Claude panel to `claude-opus-5-thinking-high`.

3.8 Descriptive taxonomy

For each claim the taxonomy records decisiveness (Supported or Refuted versus Ambiguous or Unresolved) and FEVER agreement at every stage, the first gold-agreeing stage, title and Stage C resolutions, reversals, and losses of decisiveness. Each claim then receives exactly one category, in priority order: *final nondecisive* (Stage C consensus Ambiguous or Unresolved); *decisive judge-
FEVER disagreement* (decisive Stage C verdict opposite to FEVER); *nonmonotonic judge path* (final gold agreement after a reversal or loss of decisiveness); *sentence-only gold-agreeing* (gold-agreeing from Stage A onward); *title-disclosure gold-agreeing* (first gold agreement at Stage B); and, for first gold agreement at Stage C, *additional-evidence gold-agreeing* when Stage C adds sentence text and *structured-disclosure gold-agreeing* otherwise. These categories describe where a blinded panel’s verdict stands; they are not causal mechanisms inside GPT. A decisive disagreement with FEVER is not evidence that the FEVER label is wrong, and sentence-only gold agreement is compatible with, but does not establish, a failure by GPT to use the supplied evidence. The taxonomy was specified post hoc, after inspection of the historical results, and awaits coauthor review.

3.9 Evidence audit and Stage C correction

The historical Stage C packets took each evidence line’s last tab-separated field, often a hyperlink target, instead of the sentence text. Chosen-set text differed from the canonical sentences for 936 of 1,061 cohort claims (190 of the 226 regressions); ten further claims lacked structured text because decomposed-Unicode titles were missing from a cache. Exact selected evidence sets were recovered by recomputing the original set identifiers, and all selected and alternative sentences were rebuilt from local archival FEVER data. Stage A inputs equal the GPT evidence prompts and Stage B inputs equal the historical records exactly; their frozen judgments are retained under the documented fresh-context-per-stage protocol, and all Stage C judgments are new. Consensus paths therefore combine historical Stage A/B judgments with new Stage C judgments.

A post hoc correction protocol, frozen on 2026-09-27 before any new judgment, fixed the repaired packets, the taxonomy, the statistical analyses, and the execution rules, including acceptance of the first schema-valid response for each job. Seventy required jobs (55 Grok, 15 Claude; 6,430 new Stage C judgments) all delivered schema-valid outputs.

Post-judgment amendment. Execution failures meant that some jobs needed more attempts than the protocol allowed and that its execution-record rules could not be met as written. Retries never depended on model outputs. After all outputs existed and had been inspected, we adopted a dated amendment (2026-09-28) that relaxes those execution rules and specifies that each job’s

Table 1: Paired behavioral results on 10,000 balanced FEVER Supported/Refuted claims (40,000 predictions). Accuracies and transition counts are from the verified Analysis v2 tables. Unique-claim regression counts treat a claim as one unit even if both models regress.

Model	Claim-only	+Gold ev.	Δ (pp)	Regressions	Rescues
GPT-5.4	88.69%	96.04%	+7.35	115 (1.15%)	850 (8.5%)
GPT-5.4-mini	84.67%	95.54%	+10.87	151 (1.51%)	1,238 (12.4%)
Unique claims that regress for at least one model: 226 (40 both; 186 exactly one).					

analysed output is the output its execution procedure accepted. For three jobs (two Claude, one Grok) an earlier schema-valid output exists, so we repeat every analysis with those first valid outputs substituted (Appendix A). The amendment is not preregistered; under the original rules the protocol’s primary analysis remains incomplete. Packets, models, taxonomy, statistics and analysis code are unchanged, and the protocol’s follow-up panels (a newly selected residual Claude panel and a Stage C resolver) were not run. Execution records and provenance evidence are documented in the released artifact.

3.10 Supporting cross-family data

GLM-5.3 contributes a completed Stage A panel compared to frozen Cursor Stage A ($n = 1,060$); Stage A inputs are unaffected by the audit. A historical Claude Stage C test on 231 Grok left-overs used the defective Stage C packets and a residual set selected with defective Grok Stage C judgments; it is reported only as a historical observation (Appendix C).

3.11 Statistical analysis

Paired accuracy uses McNemar tests; regression and rescue rates use bootstrap confidence intervals. Descriptive adjudication rates use Wilson 95% intervals. Cross-family comparisons on the paired 226 claims use exact McNemar tests with Benjamini–Hochberg correction across 14 tests; shared-versus-model-specific nondecisiveness uses odds ratios in a separate two-test family. Cross-family taxonomy agreement is reported as exact match and Cohen’s κ . All adjudication analyses are post hoc.

4 Results

4.1 Gold evidence improves aggregate accuracy

Table 1 answers RQ1. On the same 10,000 claims, designated gold evidence raises GPT-5.4 accuracy from 88.69% to 96.04% (+7.35 pp) and GPT-5.4-mini from 84.67% to 95.54% (+10.87 pp). Paired McNemar tests are large (GPT-5.4 $\chi^2 = 558.3$, $p \approx 2 \times 10^{-123}$; mini $\chi^2 = 849.1$, $p \approx 10^{-186}$). Figure 1 shows the four-way transition structure behind those averages. Evidence also increases cross-model prediction agreement, from 89.3% ($\kappa = 0.786$) without evidence to 97.2% ($\kappa = 0.944$) with evidence.

Table 2: Adjudication outcomes for the 226 unique regressions on repaired Stage C inputs, as counts with percentages of 226. Δ is Claude minus Grok in percentage points with exact McNemar tests, Benjamini–Hochberg adjusted over 14 cross-family tests. “FV” gives the Claude count under the first-valid sensitivity selection (= unchanged); Grok regression-cohort values are identical under both selections. Categories describe blinded judge-panel verdicts; they are not causes of GPT errors or corrected FEVER labels.

Category	Grok	Claude	FV	Δ (pp)	p_{BH}
Sentence-only gold-agreeing	77 (34.1%)	106 (46.9%)	105	+12.8	4.6×10^{-7}
Title-disclosure gold-agreeing	10 (4.4%)	3 (1.3%)	=	−3.1	0.13
Structured-disclosure gold-agreeing	2 (0.9%)	3 (1.3%)	6	+0.4	1.00
Additional-evidence gold-agreeing	17 (7.5%)	14 (6.2%)	=	−1.3	0.58
Nonmonotonic judge path	5 (2.2%)	5 (2.2%)	=	0.0	1.00
Decisive judge–FEVER disagreement	44 (19.5%)	54 (23.9%)	=	+4.4	0.015
Final nondecisive	71 (31.4%)	41 (18.1%)	39	−13.3	7.2×10^{-8}

Exact per-claim category agreement: 172/226, 76.1% (Wilson 70.1–81.2), $\kappa = 0.673$; first-valid 168/226, 74.3%, $\kappa = 0.651$

Table 3: Stage-wise nondecisiveness and resolution for the 226 regressions (percentages of 226). Stage A and B verdicts are retained historical judgments on unchanged inputs; Stage C verdicts are new judgments on repaired inputs. Δ is Claude minus Grok with exact McNemar tests, Benjamini–Hochberg adjusted over 14 tests; losses of decisiveness are counts without a test. Brackets give the first-valid sensitivity selection where it differs.

Measure	Grok	Claude	Δ (pp)	p_{BH}
Nondecisive at Stage A	103 (45.6%)	66 (29.2%)	−16.4	1.1×10^{-8}
Nondecisive at Stage B	95 (42.0%)	64 (28.3%)	−13.7	9.4×10^{-7} [1.2×10^{-6}]
Nondecisive at Stage C	71 (31.4%)	41 (18.1%) [39]	−13.3 [−14.2]	7.2×10^{-8} [9.1×10^{-8}]
Resolved by titles (A→B)	16 (7.1%)	9 (4.0%)	−3.1	0.25
Resolved at Stage C (B→C)	27 (11.9%)	24 (10.6%) [27]	−1.3 [0.0]	0.93 [1.00]
Lost decisiveness (A→B / B→C)	8 / 3	7 / 1 [7 / 2]		
Switched to opposite label	0	0	0.0	1.0

4.2 Evidence can induce correct-to-wrong regressions

The same paired design produces 115 GPT-5.4 regressions (1.15%, bootstrap 95% CI [0.95, 1.36]) and 151 mini regressions (1.51%, CI [1.28, 1.75]). Rescues are more common: 850 (8.5%) and 1,238 (12.4%). Among claim-only errors, evidence corrects 75.2% (GPT-5.4) and 80.8% (mini).

Unique-claim counts, not model-observations, are the regression cohort for later adjudication: 226 claims regress for at least one model, 40 for both, and 186 for exactly one (75 GPT-5.4-only, 111 mini-only). A paired McNemar test on discordant claims finds a modest model difference in regression incidence (115 vs. 151, $\chi^2 = 6.59$, $p = 0.010$). 87.3% of claims share the same transition class across models. The answer to RQ2 is therefore that regressions are rare relative to the 10,000 claims and only partly shared between the two models, but numerous enough to study as a structured minority. These behavioral results come from the original 40,000 predictions and are unaffected by the adjudication audit; they are observed outcomes of the original prompts, not evidence of repeat-run stability.

Table 4: Final Stage C nondecisiveness among regressions shared by both GPT models versus regressions specific to one model. Odds ratios compare shared with model-specific regressions (two-test Benjamini–Hochberg family). Brackets give the Claude first-valid sensitivity selection.

Subgroup	n	Grok final nondecisive	Claude final nondecisive
Shared by both GPT models	40	18 (45.0%)	14 (35.0%)
Exactly one GPT model	186	53 (28.5%)	27 (14.5%) [25 (13.4%)]
Odds ratio (95% CI)		2.05 (1.02–4.13)	3.17 (1.47–6.83) [3.47]
p_{BH}		0.059	0.011 [0.0046]

Q1: Four-way evidence transition classes per model (n=10,000 each)




Figure 1: Paired evidence-transition classes for GPT-5.4 and GPT-5.4-mini on the same 10,000 claims: robust (correct→correct), rescue (wrong→correct), resistant (wrong→wrong) and regression (correct→wrong).

4.3 Characteristics of evidence-induced regressions

Regressions are not a random slice of the sample. A local NLI diagnostic (nli-deberta-v3-small) disagrees with the FEVER label on 74.8% of GPT-5.4 regressions and 76.8% of mini regressions, against 22.6% of GPT-5.4 robust successes (Figure 2). Resistant failures look similar (75.8%). A second, architecturally different NLI model reproduces the same ordering. The 2,462 claims flagged as weakly warranted also have lower claim-only GPT-5.4 accuracy (81.0% vs. 88.69%), so the NLI signal marks hard items rather than a quirk of one evidence-condition model.

Label asymmetry is model-dependent. GPT-5.4 regressions concentrate on Supported claims (88 vs. 27 Refuted, $p = 1.8 \times 10^{-8}$); mini shows no such split (74 vs. 77, $p = 0.87$). Surface features (length, overlap, evidence-set size) only weakly separate regressions. Predictive models that include NLI features do better than surface-only models, but rare-class PR-AUC remains low. We treat this block as characterization; it does not depend on the adjudication stage that was corrected.

4.4 Adjudication outcomes on repaired evidence

Table 2 and Figure 3 answer RQ3 for the 226 regressions. Values in this and the next subsection use the accepted Stage C outputs; the first-valid sensitivity selection is given where it differs (Sec-

QE: Local NLI / FEVER disagreement by transition class (gpt-5.4)




Figure 2: Local NLI disagreement with the FEVER label by GPT-5.4 transition class. The NLI models are diagnostics, not a relabeling of FEVER.

tion 3.9). The largest Grok categories are sentence-only gold agreement (77; 34.1%, Wilson 95% CI 28.2–40.5) and final nondecisiveness (71; 31.4%, 25.7–37.7), whose intervals overlap. Decisive verdicts that disagree with FEVER form the third group (44; 19.5%, 14.8–25.1). Title, structured and additional-evidence gold agreement together cover 29 claims (12.8%), and five claims follow a nonmonotonic path. Among regressions with a decisive Stage C consensus, the Grok verdict agrees with FEVER on 111 of 155 (71.6%).

Progressive disclosure changes a minority of verdicts (Table 3, Figure 4). Titles make 16 of the 103 Stage A-nondecisive regressions decisive, and Stage C makes 27 of the 95 Stage B-nondecisive regressions decisive. No decisive consensus switches to the opposite label at either step, although titles leave 8 previously decisive regressions nondecisive and Stage C does so for 3 more. Because Stage C adds alternative sentence text for 41 regressions, Stage C-associated changes cannot be attributed to representation alone; the 17 additional-evidence cases are separated from the two structure-only cases for that reason.

Shared regressions remain nondecisive more often than model-specific ones (Table 4): 45.0% versus 28.5% under Grok (odds ratio 2.05, $p_{\text{BH}} = 0.059$).

The answer to RQ3 is that, under Grok, regressions form a mixture of FEVER-agreeing, FEVER-disagreeing and nondecisive outcomes with no single dominant category. These outcomes do not show why GPT changed its answer, and the decisive disagreements are not evidence that FEVER is wrong.




Figure 3: Adjudication outcome categories for the 226 unique regressions on repaired Stage C inputs, by adjudicating family. Bars use the accepted outputs; black ticks mark Claude under the first-valid sensitivity selection. Grok values are identical under both selections.

4.5 Cross-family comparison

The Claude panel judges the same 226 regressions (1,130 new Stage C judgments, paired with its 2,260 retained Stage A/B judgments). Claude assigns more regressions to sentence-only gold agreement (106, 46.9%, CI 40.5–53.4; first-valid 105; +12.8 and +12.4 pp versus Grok, $p_{BH} \leq 2.1 \times 10^{-6}$) and fewer to final nondecisiveness (41, 18.1%, CI 13.7–23.7; first-valid 39; –13.3 and –14.2 pp, $p_{BH} \leq 9.1 \times 10^{-8}$). Decisive judge–FEVER disagreement is somewhat more common under Claude (54 versus 44; +4.4 pp, $p_{BH} = 0.015$). Among decisive Stage C consensus verdicts, FEVER agreement is similar across families (Claude 131/185, 70.8%; Grok 111/155, 71.6%).

Claude is more decisive at every stage (Table 3, Figure 4): nondecisive rates are 29.2%, 28.3% and 18.1% (first-valid 17.3%) versus Grok’s 45.6%, 42.0% and 31.4%. Titles resolve 9 of Claude’s 66 Stage A-nondecisive regressions and Stage C resolves 24 of its 64 Stage B-nondecisive regressions (first-valid 27). Title and Stage C resolution rates do not differ significantly between families. Neither family switches a decisive consensus to the opposite label; under Claude, 7 decisive verdicts become nondecisive with titles and 1 at Stage C (first-valid 2).

Under Claude, shared regressions are also more often nondecisive than model-specific ones (35.0% versus 14.5%; odds ratio 3.17, $p_{BH} = 0.011$; first-valid 3.47, $p_{BH} = 0.0046$).

Exact per-claim category agreement is 172/226 (76.1%, Wilson 70.1–81.2; $\kappa = 0.673$) for the accepted selection and 168/226 (74.3%, 68.3–79.6; $\kappa = 0.651$) for first-valid. The first-valid figure coincides numerically with the historical 74.3%, which measured a different, withdrawn taxonomy on defective inputs; the two are not comparable.

Selection sensitivity. Substituting the first valid outputs for the accepted ones changes Stage C consensus for six of the 226 Claude claims and for one of the 1,060 Grok claims. The Grok change moves a rescue-cohort claim from Ambiguous to Unresolved, so no Grok category and no regression




Figure 4: Nondecisive (Ambiguous or Unresolved) five-judge consensus on the 226 regressions at each disclosure stage, with Wilson 95% intervals. Stages A and B are retained historical judgments on unchanged inputs; Stage C is new judgments on repaired inputs. The open diamond is Claude Stage C under the first-valid sensitivity selection.

result changes. All 14 cross-family tests and both odds-ratio tests keep their Benjamini–Hochberg decisions, and no cross-family difference moves by more than 1.3 pp. Across all 12 combinations of completed runs for the two affected Claude jobs, zero to six claims change consensus relative to the accepted outputs and Claude Stage C Ambiguous counts range from 37 to 42 of 226.

4.6 Supporting Stage A check

GLM-5.3 Stage A, compared with frozen Cursor Stage A on 1,060 items, yields modal verdict agreement 0.862 and consensus agreement 0.831. GLM is less ambiguity-conservative (ambiguous/unresolved rate 0.223 vs. Cursor 0.291). This is Stage A robustness only; Stage A inputs were unaffected by the audit.

4.7 What depends on the judge family

Several descriptions hold under both panels. Every category occurs under both families. At each disclosure step, both panels resolve a minority of the regressions that were nondecisive before it and leave a few decisive ones nondecisive, and neither ever switches a decisive consensus to the opposite label. Under both, roughly a fifth to a quarter of the regressions end in a decisive verdict opposite to FEVER, decisive verdicts agree with FEVER at similar rates, and shared regressions are more often nondecisive than model-specific ones, although that difference is significant under Claude only.

Other descriptions depend on which family judges. Claude is less often nondecisive at every stage, by 13–16 pp, and places a larger share of regressions in sentence-only gold agreement, so the two panels disagree about which category is largest. At the level of individual claims, the panels assign the same category to roughly three regressions in four ($\kappa = 0.65$ – 0.67). The answer to RQ4 is therefore that the qualitative mixture of outcomes is shared while its proportions depend on the judge family. Because both panels are models scoring the same packets, the design cannot

determine whether Claude is better calibrated than Grok or more willing to reach a verdict on weakly supported claims.

5 Discussion

5.1 Evidence helps globally but can hurt locally

The takeaway from these findings is not that gold evidence fails. It succeeds on average: far more claims change from incorrect to correct when evidence is added (850 for GPT-5.4 and 1,238 for GPT-5.4-mini) than change from correct to incorrect. Aggregate accuracy is nevertheless the wrong unit for evaluating reliability, because it nets these movements against each other. For GPT-5.4, 115 of the 10,000 claims (1.15%) changed from correct without evidence to wrong with it; for GPT-5.4-mini, 151 (1.51%) did; and 226 distinct claims regressed for at least one model. In each case, adding correct evidence reversed a decision that was already correct, and condition-level accuracy alone conceals every one of these reversals. Evaluation reports that publish only accuracy per condition will therefore miss them.

5.2 What adjudication can and cannot say about regressions

The Grok and Claude panels reach the FEVER label from the evidence sentences alone, and keep it as titles and structured evidence are added, for 34.1% and 46.9% of the 226 regressions, respectively. On those claims, the designated evidence was sufficient for a separate model panel to reach the FEVER label, yet GPT moved away from that label after receiving the same evidence. This pattern is compatible with a failure by GPT to use the evidence, and it is consistent with prior work showing that models can ignore, mishandle, or be distracted by supplied context [Longpre et al., 2021, Shi et al., 2023, Wan et al., 2024]. It does not establish that explanation. The original runs stored predicted labels, not rationales or confidence scores that would show why GPT changed its answer, and the judges are different models reading the evidence in a different prompt. A further 19.5% (Grok) to 23.9% (Claude) of regressions end in a decisive verdict opposite to FEVER, and 18.1% (Claude) to 31.4% (Grok) end nondecisive. These outcomes describe how the panels read the evidence. A decisive verdict opposite to FEVER is disagreement with the benchmark label, not proof that the label is wrong; a nondecisive verdict is a judge outcome, not proof that the evidence is insufficient; and neither shows that GPT failed to use evidence supporting the FEVER label. The adjudication therefore does not establish how prevalent GPT evidence-utilization failures are, or what causes them.

5.3 Nondecisiveness and disagreement with FEVER

Final nondecisiveness is common in both families. This finding is expected, given prior evidence that evidence sufficiency and interpretive ambiguity are central problems in fact verification [Atanasova et al., 2022, Glockner et al., 2024]. We do not interpret nondecisive outcomes or decisive judge–FEVER disagreements as FEVER annotation errors; judge rationales that describe a conflict between claim and evidence are observations, not labels. Shared regressions, which both GPT models got wrong after receiving evidence, are more often nondecisive under both families, although the Grok difference is not significant after correction for multiple testing.

5.4 Disclosure and representation

Titles and structured evidence change a minority of verdicts, mostly by making nondecisive verdicts decisive, and never switch a decisive consensus to the opposite label. Because Stage C adds alternative sentence text for some claims, its effect mixes added information with changed representation. Separating the two would require a structure-only arm that shows the chosen evidence set without alternatives, which we did not run. We therefore make no claim about the size or portability of a pure representation effect.

5.5 Judge-family dependence as a property of the measurement

The difference between the Claude and Grok panels is itself a finding about the measurement. Research on LLM judges already shows that evaluators have their own biases and cannot simply be exchanged for one another [Zheng et al., 2023, Liu et al., 2023, Chen et al., 2024, Huang et al., 2025]. Here the same packets, rubric and taxonomy produce a 13–16 pp difference in nondecisiveness at every stage, and the panels assign different categories to about one regression in four. Category shares therefore describe an adjudicating family as much as they describe the claims. We report both panels and do not combine them into a single estimate or treat the two values as bounds on a true prevalence.

5.6 Implications for evidence-grounded LLM evaluation

Four implications follow. First, whenever a benchmark supplies designated evidence, evaluation should compare each claim with and without it, because otherwise regressions remain invisible. Second, automated diagnostic categories need a judge from a second model family before their shares are treated as stable, and decisiveness should be reported separately from agreement with the benchmark label. Third, adjudication should keep nondecisiveness as an outcome in its own right rather than forcing each claim into Supported or Refuted. Fourth, adjudication inputs deserve the same verification as model inputs: a field-extraction error silently corrupted the structured stage of our original adjudication. The natural next measurement is human review of the 54–58 regressions on which the Grok and Claude panels assign different categories, rather than another automated judge family.

6 Limitations

All adjudication is automated silver labeling. The Grok, GLM and Claude judgments are model outputs, not expert human annotation, and they create no new gold standard; no adjudication outcome was reviewed by humans. The adjudication can therefore describe how model judges read the evidence, but not how people would read it. Agreement within a panel also overstates how far independent human raters would agree, because the five judges are isolated runs of one locked model rather than independent readers. Across panels, the category shares differ by judge family, so we report both complete panels rather than a single estimate. Coverage is also uneven: the Claude panel judges only the 226 regressions, not the 1,060-claim judged cohort, and the GLM panel covers Stage A only.

FEVER labels remain the reference standard for scoring GPT, and the experiment covers a narrow slice of FEVER. Its claims are 2017-era Wikipedia claims restricted to the Supported and Refuted labels, and the evidence is designated rather than retrieved. The results therefore say nothing directly about NOT ENOUGH INFO claims, about open-web retrieval, where retrieval errors would add to the failures studied here, or about claims from other sources and periods. Both evaluated systems are GPT-5.4 variants, so the overlap between their transitions may partly reflect errors shared within one model family. Disagreement between an NLI model or a judge panel and FEVER marks candidate ambiguity; it does not show that an annotation is wrong.

The design locates where a judge panel’s verdict changes, not what happened inside GPT. Progressive disclosure observes judges, not GPT-internal mechanisms, and the original GPT runs stored predicted labels only, so no analysis in this paper can recover why a particular prediction changed.

The adjudication itself was corrected after the fact. The historical Stage C packets displayed corrupted sentence text, and the historical taxonomy conflated disagreement with FEVER with evidence utilization. The repaired Stage C and the descriptive taxonomy were specified post hoc, after the historical results and defects had been inspected; the correction protocol is not a preregistration, and the taxonomy awaits coauthor review. Stage C also adds alternative evidence text for some claims, so changes associated with Stage C are not a pure manipulation of representation. Finally, each consensus path combines retained historical Stage A and B judgments with new Stage C judgments made on a different occasion. Treating those judgments as one path relies on the historical protocol’s fresh context for every judge, batch and stage, which could not be verified independently because raw historical session transcripts were unavailable; if that isolation failed, some stage-to-stage changes could reflect earlier context rather than the newly disclosed evidence. Historical execution events in the original Claude panel (a quarantined rerun after a blinding leak, a paused and resumed stage, and a duplicate-batch tie-break) affect only its retained Stage A/B judgments and are documented in Appendix B.

The Stage C analysis also rests on an amendment adopted after the judgments and their provisional results had been seen. The amendment relaxed the protocol’s execution rules and selected the accepted outputs where earlier valid outputs existed, so under the original rules the protocol’s primary analysis is incomplete and the amended analysis is not a pre-specified result. The amendment governs which outputs are admitted, not how they are analysed, and the first-valid alternative changes six Claude consensus labels without changing any significance decision. Execution provenance rests on operator records and hashed logs rather than provider attestation. The protocol’s follow-up panels (a newly selected residual Claude panel, a Stage C resolver, and an optional fixed historical residual panel) were not run.

7 Conclusion

Designated gold evidence improves GPT-5.4 and GPT-5.4-mini fact verification in aggregate but still induces 226 unique correct-to-wrong regressions. These behavioral results are unchanged by the adjudication audit. Re-adjudicated on repaired structured evidence by two model families, the regressions form a mixture of FEVER-agreeing, FEVER-disagreeing and nondecisive panel outcomes; disclosure of titles and structured evidence resolves a minority of the nondecisive cases and never switches a decisive verdict to the opposite label. The mixture is common to both

families, but its proportions are not: Claude is 13–16 pp more decisive at every stage and places more regressions in sentence-only gold agreement, and the two families assign the same category to about three regressions in four. The historical three-way failure-mode shares and their cross-family ordering are withdrawn, and the adjudication does not identify why GPT changed its answers. Paired no-evidence versus gold-evidence evaluation, a second judge family, and verified adjudication inputs are what it takes to see these failures and describe them without overstating them.

Data and code

The behavioral record is tag `evidex-artifact-v1` at commit `29102c3`; its headline verifiers and test suites are unchanged. The evidence audit, repaired Stage C packets, frozen correction protocol (2026-09-27), dated Stage C amendment (2026-09-28), all execution records including failed attempts, and the amended and sensitivity analyses are on branch `codex/methodology-correction-v2`. Project code is MIT-licensed; FEVER and Wikipedia-derived files are third-party material.

References

- Rami Aly, Zhijiang Guo, Michael Schlichtkrull, James Thorne, Andreas Vlachos, Christos Christodoulopoulos, Oana Cocarascu, and Arpit Mittal. FEVEROUS: Fact extraction and VERification over unstructured and structured information. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, volume 1, 2021. URL <https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/68d30a9594728bc39aa24be94b319d21-Abstract-round1.html>.
- Pepa Atanasova, Jakob Grue Simonsen, Christina Lioma, and Isabelle Augenstein. Fact checking with insufficient evidence. *Transactions of the Association for Computational Linguistics*, 10:746–763, 2022. doi: 10.1162/tacl_a.00486. URL <https://aclanthology.org/2022.tacl-1.43/>.
- Guiming Hardy Chen, Shunian Chen, Ziche Liu, Feng Jiang, and Benyou Wang. Humans or LLMs as the judge? a study on judgement bias. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8301–8327, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.474. URL <https://aclanthology.org/2024.emnlp-main.474/>.
- Max Glockner, Ieva Staliūnaitė, James Thorne, Gisela Vallejo, Andreas Vlachos, and Iryna Gurevych. AmbiFC: Fact-checking ambiguous claims with evidence. *Transactions of the Association for Computational Linguistics*, 12:1–18, 2024. doi: 10.1162/tacl_a.00629. URL <https://aclanthology.org/2024.tacl-1.1/>.
- Zhijiang Guo, Michael Schlichtkrull, and Andreas Vlachos. A survey on automated fact-checking. *Transactions of the Association for Computational Linguistics*, 10:178–206, 2022. doi: 10.1162/tacl_a.00454. URL <https://aclanthology.org/2022.tacl-1.11/>.
- Hui Huang, Xingyuan Bu, Hongli Zhou, Yingqi Qu, Jing Liu, Muyun Yang, Bing Xu, and Tiejun Zhao. An empirical study of LLM-as-a-judge for LLM evaluation: Fine-tuned judge model is

- not a general substitute for GPT-4. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 5880–5895, Vienna, Austria, July 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.findings-acl.306. URL <https://aclanthology.org/2025.findings-acl.306/>.
- Yichen Jiang, Shikha Bordia, Zheng Zhong, Charles Dognin, Maneesh Singh, and Mohit Bansal. HoVer: A dataset for many-hop fact extraction and claim verification. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3441–3460, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.309. URL <https://aclanthology.org/2020.findings-emnlp.309/>.
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024. doi: 10.1162/tacl.a_00638. URL <https://aclanthology.org/2024.tacl-1.9/>.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-eval: NLG evaluation using Gpt-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2511–2522, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.153. URL <https://aclanthology.org/2023.emnlp-main.153/>.
- Shayne Longpre, Kartik Perisetla, Anthony Chen, Nikhil Ramesh, Chris DuBois, and Sameer Singh. Entity-based knowledge conflicts in question answering. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7052–7063, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-main.565. URL <https://aclanthology.org/2021.emnlp-main.565/>.
- Mamta and Oana Cocarascu. FactEval: Evaluating the robustness of fact verification systems in the era of large language models. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 10647–10660, Albuquerque, New Mexico, April 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.naacl-long.534. URL <https://aclanthology.org/2025.naacl-long.534/>.
- Tal Schuster, Darsh Shah, Yun Jie Serene Yeo, Daniel Roberto Filizzola Ortiz, Enrico Santus, and Regina Barzilay. Towards debiasing fact verification models. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3419–3425, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1341. URL <https://aclanthology.org/D19-1341/>.
- Tal Schuster, Adam Fisch, and Regina Barzilay. Get your vitamin C! robust fact verification with contrastive evidence. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 624–643, Online, June 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.naacl-main.52. URL <https://aclanthology.org/2021.naacl-main.52/>.

- Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Schärli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 31210–31227. PMLR, 2023. URL <https://proceedings.mlr.press/v202/shi23a.html>.
- Lin Shi, Chiyu Ma, Wenhua Liang, Xingjian Diao, Weicheng Ma, and Soroush Vosoughi. Judging the judges: A systematic study of position bias in LLM-as-a-judge. In *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics*, pages 292–314, Mumbai, India, December 2025. The Asian Federation of Natural Language Processing and The Association for Computational Linguistics. doi: 10.18653/v1/2025.ijcnlp-long.18. URL <https://aclanthology.org/2025.ijcnlp-long.18/>.
- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. FEVER: a large-scale dataset for fact extraction and VERification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 809–819, New Orleans, Louisiana, June 2018a. Association for Computational Linguistics. doi: 10.18653/v1/N18-1074. URL <https://aclanthology.org/N18-1074/>.
- James Thorne, Andreas Vlachos, Oana Cocarascu, Christos Christodoulopoulos, and Arpit Mittal. The fact extraction and VERification (FEVER) shared task. In *Proceedings of the First Workshop on Fact Extraction and VERification (FEVER)*, pages 1–9, Brussels, Belgium, November 2018b. Association for Computational Linguistics. doi: 10.18653/v1/W18-5501. URL <https://aclanthology.org/W18-5501/>.
- David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. Fact or fiction: Verifying scientific claims. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7534–7550, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.609. URL <https://aclanthology.org/2020.emnlp-main.609/>.
- Alexander Wan, Eric Wallace, and Dan Klein. What evidence do language models find convincing? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7468–7484, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.403. URL <https://aclanthology.org/2024.acl-long.403/>.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html. Datasets and Benchmarks Track.

A Stage C amendment and selection sensitivity

The correction protocol (frozen 2026-09-27, before any new judgment) required the first schema-valid response of each job to be analysed. The dated amendment (2026-09-28), written after all 70 outputs existed and had been inspected, analyses each job’s accepted output instead and relaxes the protocol’s attempt and execution-record rules. Every other provenance check (locked model routing, fresh context, input and output hashes, and raw output identical to the delivered output) is applied unchanged, as are the packets, taxonomy, statistics and analysis code.

Three jobs have an earlier schema-valid output than the accepted one: two Claude jobs, with two and three earlier complete outputs, and one Grok job, whose first run wrote a complete output before ending with an error. The first-valid sensitivity selection substitutes the earliest of these outputs in all three jobs.

Quantity (of 226 regressions)	Accepted	First-valid
Claude Stage C nondecisive	41 (18.1%)	39 (17.3%)
Claude sentence-only gold-agreeing	106	105
Claude structured-disclosure gold-agreeing	3	6
Claude Stage C consensus labels changed	—	6
Grok Stage C labels changed (of 1,060 claims)	—	1, not a regression
Exact Grok–Claude category agreement	172 (76.1%), $\kappa = 0.673$	168 (74.3%), $\kappa = 0.651$
Claude shared-vs-specific odds ratio (p_{BH})	3.17 (0.011)	3.47 (0.0046)
Benjamini–Hochberg decisions changed (16 tests)	—	0

Across all 12 combinations of the completed runs for the two Claude jobs, Stage C consensus differs from the accepted selection on zero to six claims, and the Claude Stage C Ambiguous count ranges from 37 to 42.

B Historical execution events in the Claude full panel

The historical Claude full panel reused the locked model `claude-opus-5-thinking-high`, the shared rubric, and the blinded Stage A/B/C evidence files; its Stage A/B judgments are retained in this draft and its historical Stage C judgments are not used. During the first Stage A batch-01 wave, judge-facing file paths included the panel directory name and therefore the cohort-selection criterion. Before inspecting the affected judgments, the five delivered files ($5 \times 77 = 385$ valid records) were quarantined, judge I/O moved to a neutral directory, and that batch rerun under the unchanged model, rubric, and evidence packets. No quarantined verdict was read. Stage B was paused after 601 of 1,130 judgments and resumed on the same Claude model; already-valid judgments were kept. Recovery overlap caused three judge-batches to be produced twice; a first-ingested-wins rule, applied without reading either copy’s labels, retained one file per (stage, judge, item).

C Withdrawn historical diagnostic results

The following results appeared in the historical draft (`paper/`, preserved unchanged). They were computed with the legacy three-way taxonomy on Stage C packets that displayed corrupted sentence text, and they are withdrawn as estimates. They are listed only so readers can trace the correction.

- Grok three-way split of the 226 regressions: 50.0% evidence-utilization failure, 11.9% representation-sensitive, 38.1% residual ambiguity (113/27/86). Of the 113 legacy utilization cases, 36 had a decisive Stage C verdict opposite to FEVER.
- Claude three-way split: 65.9% / 5.3% / 28.8% (149/12/65); 46 of the 149 legacy utilization cases disagreed with FEVER at Stage C.
- “Utilization failure is largest in both panels” and the claimed replication of the category ordering; the 50–66% range for utilization failure.
- Legacy taxonomy agreement 74.3% (168/226), $\kappa = 0.537$, and the four-item representation-sensitive overlap.
- Historical Stage C nondecisive rates (Grok 38.1%, Claude 28.8%) and the Stage C-dependent McNemar test.
- The historical Claude residual test (134/231, 58.0%, persistently ambiguous), whose inputs and residual selection both depended on defective Stage C.

Historical figures of the category split, taxonomy confusion, Stage C ambiguity and Grok resolution flow are not reproduced in this draft.

D Denominator checklist

Quantity	<i>n</i>
FEVER claims in the paired experiment	10,000
GPT predictions	40,000
Diagnostic-cohort claims constructed	1,061
Judged diagnostic-cohort items	1,060
Unique regressions	226
New Stage C judgments, Grok ($1,060 \times 5$)	5,300
New Stage C judgments, Claude (226×5)	1,130
Required correction jobs delivered	70
Retained historical Claude Stage A/B judgments ($226 \times 5 \times 2$)	2,260
Historical Claude residual cohort (Grok leftovers; withdrawn test)	231
Optional fixed historical-residual judgments (231×5 ; not run)	1,155